

Task 1: Familiarize Yourself with the Dataset

Import all the datasets

```
In [90]: import pandas as pd

sessions_df = pd.read_csv("data/sessions.csv")
streamers_df = pd.read_csv("data/streamers.csv")
streams_df = pd.read_csv("data/streams.csv")
viewers_df = pd.read_csv("data/viewers.csv")
```

Read the head of all the datasets to understand column format

```
In [91]: sessions_df.head()
```

```
Out[91]:
```

	session_id	viewer_id	streamer_id	stream_id	started_at	ended_at	duration_mins	chat_messages_s
0	SES000001	VWR00938	STR0055	SM00860	2024-01-18 19:16:00	2024-01-18 20:28:00	72	
1	SES000002	VWR00804	STR0034	SM00542	2024-01-04 17:11:00	2024-01-04 19:41:00	150	
2	SES000003	VWR00704	STR0008	SM00124	2024-02-15 15:11:00	2024-02-15 16:53:00	102	
3	SES000004	VWR00926	STR0024	SM00373	2024-01-18 19:59:00	2024-01-18 20:56:00	57	
4	SES000005	VWR00883	STR0008	SM00120	2024-03-19 20:08:00	2024-03-19 20:22:00	14	

```
In [92]: sessions_df.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 22825 entries, 0 to 22824
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   session_id            22825 non-null  str
1   viewer_id             22825 non-null  str
2   streamer_id           22765 non-null  str
3   stream_id             22825 non-null  str
4   started_at            22825 non-null  str
5   ended_at              22825 non-null  str
6   duration_mins          22825 non-null  int64
7   chat_messages_sent    22825 non-null  int64
8   bits_cheered          22825 non-null  int64
9   followed_during       22825 non-null  bool
10  subscribed_during      22825 non-null  bool
dtypes: bool(2), int64(3), str(6)
memory usage: 1.6 MB
```

```
In [93]: streamers_df.head()
```

Out[93]:

	streamer_id	streamer_name	category	language	partner_status	total_followers	avg_concurrent_viewer
0	STR0001	xQc	Just Chatting	English	True	11500000	13870
1	STR0002	Pokimane	Just Chatting	English	True	-25551	7166
2	STR0003	HasanAbi	Just Chatting	English	True	2600000	3422
3	STR0004	Amouranth	Just Chatting	English	True	6200000	6860
4	STR0005	Mizkif	Just Chatting	English	True	2100000	677

In [94]: streamers_df.info()

<class 'pandas.DataFrame'>
RangeIndex: 76 entries, 0 to 75
Data columns (total 7 columns):
Column Non-Null Count Dtype
--- -
0 streamer_id 76 non-null str
1 streamer_name 76 non-null str
2 category 76 non-null str
3 language 76 non-null str
4 partner_status 76 non-null bool
5 total_followers 76 non-null int64
6 avg_concurrent_viewers 76 non-null int64
dtypes: bool(1), int64(2), str(4)
memory usage: 3.8 KB

In [95]: streams_df.head()

Out[95]:

	stream_id	streamer_id	started_at	ended_at	stream_duration_hrs	category	peak_viewers	title_has_
0	SM00001	STR0001	2024-02-14 18:00:00	2024-02-14 21:00:00	3.0	Just Chatting	108989.0	
1	SM00002	STR0001	2024-01-28 18:00:00	2024-01-28 21:00:00	3.0	Just Chatting	51635.0	
2	SM00003	STR0001	2024-01-31 09:00:00	2024-01-31 16:00:00	7.0	Just Chatting	212127.0	
3	SM00004	STR0001	2024-03-03 14:00:00	2024-03-03 21:00:00	7.0	Just Chatting	186745.0	
4	SM00005	STR0001	2024-01-24 18:00:00	2024-01-24 21:00:00	3.0	Just Chatting	87723.0	

In [96]: streams_df.info()

```
<class 'pandas.DataFrame'>
RangeIndex: 1199 entries, 0 to 1198
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   stream_id             1199 non-null  str
1   streamer_id           1199 non-null  str
2   started_at            1199 non-null  str
3   ended_at              1199 non-null  str
4   stream_duration_hrs   1199 non-null  float64
5   category              1149 non-null  str
6   peak_viewers          1119 non-null  float64
7   title_has_hype_word   1199 non-null  bool
8   was_raided            1199 non-null  bool
dtypes: bool(2), float64(2), str(5)
memory usage: 68.0 KB
```

```
In [97]: viewers_df.head()
```

```
Out[97]:
```

	viewer_id	age_group	country	account_age_days	subscription_tier	preferred_category
0	VWR00001	18-24	AR	355	free	IRL
1	VWR00002	25-34	CA	1670	free	Creative
2	VWR00003	35-44	US	1108	free	Just Chatting
3	VWR00004	45+	NaN	822	tier1	FPS
4	VWR00005	25-34	DE	2226	free	Just Chatting

```
In [98]: viewers_df.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   viewer_id             1000 non-null  str
1   age_group             1000 non-null  str
2   country               920 non-null   str
3   account_age_days      1000 non-null  int64
4   subscription_tier     1000 non-null  str
5   preferred_category    1000 non-null  str
dtypes: int64(1), str(5)
memory usage: 47.0 KB
```

Covert dates/times into pandas date/time objects to manipulate easier

```
In [99]: # sessions.csv
sessions_df['started_at'] = pd.to_datetime(sessions_df['started_at'], format='%Y-%m-%d %H:%M:%S')
sessions_df['ended_at'] = pd.to_datetime(sessions_df['ended_at'], format='%Y-%m-%d %H:%M:%S')

# streams.csv
streams_df['started_at'] = pd.to_datetime(streams_df['started_at'], format='%Y-%m-%d %H:%M:%S')
streams_df['ended_at'] = pd.to_datetime(streams_df['ended_at'], format='%Y-%m-%d %H:%M:%S')

In [100... # check that it worked
sessions_start_date = sessions_df['started_at']
streams_start_date = streams_df['started_at']
print(sessions_start_date)
```

```
print()
print(streams_start_date)
```

```
0      2024-01-18 19:16:00
1      2024-01-04 17:11:00
2      2024-02-15 15:11:00
3      2024-01-18 19:59:00
4      2024-03-19 20:08:00
...
22820   2024-01-16 18:51:00
22821   2024-01-08 20:28:00
22822   2024-03-26 04:15:00
22823   2024-01-11 20:02:00
22824   2024-01-13 11:23:00
Name: started_at, Length: 22825, dtype: datetime64[us]

0      2024-02-14 18:00:00
1      2024-01-28 18:00:00
2      2024-01-31 09:00:00
3      2024-03-03 14:00:00
4      2024-01-24 18:00:00
...
1194    2024-01-31 19:00:00
1195    2024-03-20 19:00:00
1196    2024-01-01 14:00:00
1197    2024-01-02 20:00:00
1198    2024-02-19 08:00:00
Name: started_at, Length: 1199, dtype: datetime64[us]
```

Task 1 Summary Report

I loaded the head and the general info of each dataset to quickly summarize whats in each set. When there was a date or time value, I converted it to a pandas datetime object.

1. sessions.csv = sessions_df

- total 11 columns
- range: 22825 entries, 0 to 22824
- dtypes: bool(2), datetime64us, int64(3), str(4)
- 'streamer_id' has 100 null entries

-
- *session_id, viewer_id, streamer_id, stream_id* = str
 - *started_at, ended_at* = datetime64[us]
 - *duration_mins, chat_messages_sent, bits_cheered* = int64
 - *followed_during, subscribed_during* = bool

2. streamers.csv = streamers_df

- total 7 columns
 - range: 76 entries, 0 to 75
 - dtypes: bool(1), int64(2), str(4)
 - 'language' has 8 null entries
-

- *streamer_id, streamer_name, category, language* = str
- *total_followers, avg_concurrent_viewers* = int64
- *partner_status* = bool

3. streams.csv = streams_df

- total 9 columns
- range: 1199 entries, 0 to 1198
- dtypes: bool(2), datetime64[us](2), float64(2), str(3)
- 'category' has 50 null entries

-
- *category, streamer_id, stream_id* = str
 - *started_at, ended_at* = datetime64[us]
 - *peak_viewers, stream_duration_hrs* = float64
 - *title_has_hype_word, was_raided* = bool

4. viewers.csv = viewers_df

- total 6 columns
- range: 1000 entries, 0 to 999
- dtypes: int64(1), str(5)
- 'country' has 80 null entries

-
- *age_group, viewer_id, country, subscription_tier, preferred_category* = str
 - *account_age_days* = int64

Foreign Keys

sessions, streamers, & streams: *streamer_id*

viewers & sessions: *viewer_id*

sessions & streams: *stream_id*

Task 2: Data Cleaning, Diagnostics, and Corrections

Systemic Problems

viewers.csv

- Missing values for 'country' (80)

Resolved by adding new country label 'unknown' to stay specific but not drop potentially important entries.

In [101...

```
viewers_df_resolved = viewers_df.copy()

viewers_df_resolved['country'] = viewers_df_resolved['country'].fillna('unknown')

viewers_df_resolved.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   viewer_id             1000 non-null   str
1   age_group             1000 non-null   str
2   country               1000 non-null   str
3   account_age_days      1000 non-null   int64
4   subscription_tier      1000 non-null   str
5   preferred_category     1000 non-null   str
dtypes: int64(1), str(5)
memory usage: 47.0 KB
```

sessions.csv

- Missing values for 'streamer_id' (100)

Resolved by adding new value 'unknown' to stay specific but not drop potentially important entries.

In [102...

```
sessions_df_resolved = sessions_df.copy()

sessions_df_resolved['streamer_id'] = sessions_df_resolved['streamer_id'].fillna('unknown')

sessions_df_resolved.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 22825 entries, 0 to 22824
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   session_id            22825 non-null   str
1   viewer_id             22825 non-null   str
2   streamer_id           22825 non-null   str
3   stream_id             22825 non-null   str
4   started_at            22825 non-null   datetime64[us]
5   ended_at              22825 non-null   datetime64[us]
6   duration_mins         22825 non-null   int64
7   chat_messages_sent    22825 non-null   int64
8   bits_cheered          22825 non-null   int64
9   followed_during       22825 non-null   bool
10  subscribed_during     22825 non-null   bool
dtypes: bool(2), datetime64[us](2), int64(3), str(4)
memory usage: 1.6 MB
```

- Some entries for 'started_at' and 'ended_at' are swapped
 - For instance, some streams end at an earlier date than when it started.

Found by Tim K. on Canvas

For this, I decided to swap the values as they seemed to be an input glitch and not any other kind of error. Once swapped, the values make sense.

In [103...

```
# create mask for sessions that end earlier than they start and swap them
swap_mask = sessions_df["started_at"] > sessions_df["ended_at"]

#check which ones are the broken entries
sessions_df.loc[swap_mask].head()
```

Out[103...

	session_id	viewer_id	streamer_id	stream_id	started_at	ended_at	duration_mins	chat_message
341	SES000342	VWR00878	STR0010	SM00147	2024-03-19 13:53:00	2024-03-19 12:15:00	98	
722	SES000723	VWR00859	STR0020	SM00311	2024-03-15 14:53:00	2024-03-15 11:57:00	176	
1011	SES001012	VWR00722	STR0076	SM01199	2024-02-19 08:36:00	2024-02-19 08:18:00	18	
1718	SES001719	VWR00916	STR0073	SM01160	2024-02-13 19:38:00	2024-02-13 19:36:00	2	
1897	SES001898	VWR00947	STR0020	SM00303	2024-01-01 11:18:00	2024-01-01 09:14:00	124	

In [104...

```
#swap the values for ended_at and started_at
sessions_df_resolved.loc[swap_mask, ["started_at", "ended_at"]] = (
    sessions_df.loc[swap_mask, ["ended_at", "started_at"]].values
)

sessions_df_resolved.head(341)
```

Out[104...

	session_id	viewer_id	streamer_id	stream_id	started_at	ended_at	duration_mins	chat_message
0	SES000001	VWR00938	STR0055	SM00860	2024-01-18 19:16:00	2024-01-18 20:28:00	72	
1	SES000002	VWR00804	STR0034	SM00542	2024-01-04 17:11:00	2024-01-04 19:41:00	150	
2	SES000003	VWR00704	STR0008	SM00124	2024-02-15 15:11:00	2024-02-15 16:53:00	102	
3	SES000004	VWR00926	STR0024	SM00373	2024-01-18 19:59:00	2024-01-18 20:56:00	57	
4	SES000005	VWR00883	STR0008	SM00120	2024-03-19 20:08:00	2024-03-19 20:22:00	14	
...	...	...	...	...	...	...	...	...
336	SES000337	VWR00886	STR0032	SM00511	2024-03-09 22:30:00	2024-03-09 23:08:00	38	
337	SES000338	VWR00624	STR0070	SM01105	2024-03-30 23:31:00	2024-03-30 23:33:00	2	
338	SES000339	VWR00680	STR0074	SM01172	2024-03-05 17:45:00	2024-03-05 19:08:00	83	
339	SES000340	VWR00435	STR0038	SM00611	2024-02-04 12:05:00	2024-02-04 12:19:00	14	
340	SES000341	VWR00670	STR0051	SM00792	2024-02-19 12:50:00	2024-02-19 13:10:00	20	

341 rows × 11 columns

- Session records reference non-existent streamers AND streams.

Found by Tawshif Al Mehrab on Canvas

Resolved by removing these entries as they might be problematic for analysis.

I used this stackoverflow to help me combine both detections with a conditional: <https://stackoverflow.com/questions/23199796/detect-and-exclude-outliers-in-a-pandas-dataframe>

In [105...

```
#rows where stream_id is missing from streams_df but present in sessions
#sessions_streamid_outliers = sessions_df[~sessions_df["stream_id"].isin(streams_df["stream_id"])]

#rows where streamer_id is missing from streamers_df but present in sessions
#sessions_streamerid_outliers = sessions_df[~sessions_df["streamer_id"].isin(streamers_df["streamer_id"])]

#both conditions combined with either/or statement
#this contains all the rows with broken IDs in sessions_df
sessions_missingid_outliers = sessions_df[
    (~sessions_df["stream_id"].isin(streams_df["stream_id"])) |
    (~sessions_df["streamer_id"].isin(streamers_df["streamer_id"]))
]

#this only keeps the rows that aren't broken, we copy it to the new resolved df
sessions_df_resolved = sessions_df[
```



```
(sessions_df["stream_id"].isin(streams_df["stream_id"])) &
(sessions_df["streamer_id"].isin(streamers_df["streamer_id"]))
]
```

```
In [106... sessions_missingid_outliers.head()
```

	session_id	viewer_id	streamer_id	stream_id	started_at	ended_at	duration_mins	chat_message:
126	SES000127	VWR00930	STR0018	SM99031	2024-03-10 16:24:00	2024-03-10 18:08:00	104	
475	SES000476	VWR00965	STR0037	SM99030	2024-01-02 00:21:00	2024-01-02 00:24:00	3	
489	SES000490	VWR00763	NaN	SM01155	2024-03-05 10:25:00	2024-03-05 10:31:00	6	
519	SES000520	VWR00509	NaN	SM01180	2024-01-02 20:26:00	2024-01-02 21:33:00	67	
847	SES000848	VWR00470	NaN	SM00355	2024-02-05 20:25:00	2024-02-05 20:42:00	17	

```
In [107... sessions_df_resolved.info()
```

```
<class 'pandas.DataFrame'>
Index: 22730 entries, 0 to 22824
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   session_id            22730 non-null  str
1   viewer_id             22730 non-null  str
2   streamer_id           22730 non-null  str
3   stream_id             22730 non-null  str
4   started_at            22730 non-null  datetime64[us]
5   ended_at              22730 non-null  datetime64[us]
6   duration_mins         22730 non-null  int64
7   chat_messages_sent    22730 non-null  int64
8   bits_cheered          22730 non-null  int64
9   followed_during       22730 non-null  bool
10  subscribed_during     22730 non-null  bool
dtypes: bool(2), datetime64[us](2), int64(3), str(4)
memory usage: 1.8 MB
```

streamers.csv

- Some follower counts are in the negative which is not possible. This is probably an input error.

Simply turned these into positive numbers by using pandas absolute function.

```
In [108... streamers_df["total_followers"] = streamers_df["total_followers"].abs()
```

- Missing values for 'language' (8)

Manually inputted in the language in the csv by looking up the streamers with that value missing

streams.csv

- 'category' has 50 null entries

Resolved by adding new value 'unknown' to stay specific but not drop potentially important entries.

In [109...

```
streams_df_resolved = streams_df.copy()

streams_df_resolved['category'] = streams_df_resolved['category'].fillna('unknown')

streams_df_resolved.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 1199 entries, 0 to 1198
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   stream_id             1199 non-null   str
1   streamer_id           1199 non-null   str
2   started_at            1199 non-null   datetime64[us]
3   ended_at              1199 non-null   datetime64[us]
4   stream_duration_hrs   1199 non-null   float64
5   category              1199 non-null   str
6   peak_viewers          1119 non-null   float64
7   title_has_hype_word    1199 non-null   bool
8   was_raided            1199 non-null   bool
dtypes: bool(2), datetime64[us](2), float64(2), str(3)
memory usage: 68.0 KB
```

- 'peak_viewers' has 80 null entries

Resolved by replacing it with the median instead of mean due to outliers being present in the form of viral streams.

In [110...

```
streams_viewer_median = streams_df_resolved['peak_viewers'].median().round(2)

streams_df_resolved['peak_viewers'] = streams_df['peak_viewers'].fillna(streams_viewer_median)

print(f"Median used for replacement: {streams_viewer_median:.2f}")

# Check data has been fully resolved by looking at rows where peak viewers & category is missing
streams_df_resolved.head(25).tail(5)
```

Median used for replacement: 20224.00

Out[110...

	stream_id	streamer_id	started_at	ended_at	stream_duration_hrs	category	peak_viewers	title_has
20	SM00021	STR0002	2024-03-23 22:00:00	2024-03-23 23:30:00	1.5	unknown	317.0	
21	SM00022	STR0002	2024-02-19 13:00:00	2024-02-19 18:00:00	5.0	unknown	20224.0	
22	SM00023	STR0002	2024-01-29 11:00:00	2024-01-29 13:30:00	2.5	Just Chatting	300.0	
23	SM00024	STR0002	2024-03-22 11:00:00	2024-03-22 18:00:00	7.0	Just Chatting	183.0	
24	SM00025	STR0002	2024-01-05 12:00:00	2024-01-05 20:00:00	8.0	Just Chatting	463.0	

Heuristics

viewers.csv

- 'free', 'free trial', and 'none' are the same account type

Replaced with 'free'

In [162...

```
account_combine = ["free", "free trial", "none"]

viewers_df["subscription_tier"] = viewers_df_resolved["subscription_tier"].replace(account_combine)
```

- 'tier 4' and 'premium' as an account type does not actually exist. The maximum tier is 3.
- Some (5) rows contain 'TIER1' instead of 'tier1'

Found by Johann Ferrera and Zoe Mann on Canvas

replaced with 'tier3'

'TIER1' replaced with 'tier1'

In [163...

```
#replace tier4/preimum withn tier3
nonexistent_acc_types = ["tier4", "premium"]

viewers_df_resolved["subscription_tier"] = viewers_df["subscription_tier"].replace(nonexistent_acc_types)

#replace TIER1 with tier1
viewers_df_resolved["subscription_tier"] = viewers_df_resolved["subscription_tier"].replace("TIER1", "tier1")
```

- Inconsistent age data convention, for instance '18 to 24' instead of the usual '18-24'

Additionally, some inputs use an en-dash instead of a hyphen. Found by Michael Mora & Atif Mahmud on Canvas.

Replace them in their strings directly, regex=False means it will treat those specific characters as literals not conditionals

In [164...

```
# replace 'to' to - using original dataset first
viewers_df_resolved["age_group"] = viewers_df["age_group"].str.replace(" to ", "-", regex=False)

# replace endash to - in resolved one
viewers_df_resolved["age_group"] = viewers_df_resolved["age_group"].str.replace("–", "-", regex=False)

#double check, in the .csv we can see entry 110 has '35 to 44'
#print to see if its changed to a hyphen
viewers_df_resolved.head(111)
```

Out [164...

	viewer_id	age_group	country	account_age_days	subscription_tier	preferred_category
0	VWR00001	18-24	AR	355	free	IRL
1	VWR00002	25-34	CA	1670	free	Creative
2	VWR00003	35-44	US	1108	free	Just Chatting
3	VWR00004	45+	unknown	822	tier1	FPS
4	VWR00005	25-34	DE	2226	free	Just Chatting
...	...	...	...	...	...	...
106	VWR00107	35-44	US	1802	free	MOBA
107	VWR00108	13-17	AU	3437	free	Just Chatting
108	VWR00109	35-44	MX	2251	tier1	FPS
109	VWR00110	25-34	AU	1519	free	Sports
110	VWR00111	18-24	ES	1574	free	Sports

111 rows × 6 columns

sessions.csv

Some streams have the same started_at and ended_at time/date

Some bits and followers were gained during these streamers which is not possible

Decided to remove them as they are problematic entries with more than one problematic value, such as gaining bits or followers.

In [165...

```
# only keep rows where started_at and ended_at are not the same
sessions_df_resolved = sessions_df_resolved[sessions_df_resolved["started_at"] != sessions_df_res
```

Task 2 summary report

Systemic Problems

viewers.csv

- **Missing Country Values (80 entries):** Applied `.fillna('unknown')` to retain rows with missing fields while still flagging them

sessions.csv

- **Missing Streamer IDs (100 entries):** Handled the same to the missing country values using `.fillna('unknown')`
- **Swapped Timestamps:** Fixed an input glitch where stream end times occurred before start times. Used a boolean comparison mask and extracted matrix properties with `.values` to swap values.
- **Relational Outliers:** Cleaned the dataset by applying AND (`&`) operator to retain only rows whose IDs matched valid IDs in `streams_df` and `streamers_df`.

streamers.csv

- **Negative Follower Counts:** Turned negative inputs into positive ones using `.abs()` which looked like an input glitch
- **Missing Language Values (8 entries):** Manually cross-referenced and updated the missing information directly.

streams.csv

- **Missing Category Entries (50 entries):** Applied `.fillna('unknown')` to missing entries without dropping the rows.

Heuristic Adjustments

viewers.csv

- **Account Type Consolidation:** Combined multiple overlapping terms ('free', 'free trial', 'none') into a single `free` category using `.replace()`.
- **Invalid & Casing Anomalies:** Fixed non-existent tier classifications ('tier4', 'premium') by remapping them using `.replace()`. Fixed inconsistencies like `'TIER1'` by replacing them with the lowercase standard format
- **Age Group Format Alignment:** Fixed mismatched conventions (e.g., '18 to 24' and en-dashes) with `.str.replace()` to a standard hyphenated format (18-24).

sessions.csv

- **Identical Start/End Timestamps:** Removed zero-duration streaming logs containing impossible metrics (such as recorded bit/follower gains during non-existent runtimes) by using the inequality operator (`!=`).

Task 3: Descriptive Statistics and Visualizations

1. Summary statistics to understand user characteristics & age groups

This includes tiers, age groups, and individual viewer interest

Age Group of Viewers by Country

In [115...

```
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
import altair as alt
import numpy as np


# All the dataframes for reference:
# sessions_df_resolved, streams_df_resolved, viewers_df_resolved, streamers_df

sns.set_theme(style="ticks") # add ticks to the axes
```

```
f, ax = plt.subplots(figsize=(10, 10))
sns.despine(f)

sns.histplot(
    viewers_df_resolved,
    x="age_group", hue="country",
    multiple="stack", #put on top of eachother
    palette="tab10",
    edgecolor="1", #stroke 1=white
    linewidth=.5,
)
```

Out[115... <Axes: xlabel='age_group', ylabel='Count'>



Altair version of the graph using the documentation

- https://altair-viz.github.io/altair-viz-v4/gallery/horizontal_stacked_bar_chart.html
- https://altair-viz.github.io/altair-viz-v4/user_guide/customization.html?highlight=color


I first tried with Seaborn since thats what we did in the lab, but I like Altair better because its similar to what

we did in vega-lite

In [116...

```
alt.Chart(viewers_df_resolved).mark_bar().encode(
    x='count():Q',
    y='age_group:N',
    color=alt.Color('country', scale=alt.Scale(scheme='tableau20')),
).properties(
    width=700,
    height=alt.Step(30)
)
```

Out[116...



Viewer Count by Subscription Tier


Used altair methods `x=alt.x()` for this instead of keywords & tooltips

- https://altair-viz.github.io/user_guide/encodings/index.html

In [117...

```
alt.Chart(viewers_df_resolved).mark_bar(color="purple").encode(
    x=alt.X(
        "subscription_tier:N",
        title="Subscription Tier",
        sort=["free", "tier1", "tier2", "tier3"],
    ),
    y=alt.Y("count():Q", title="Number of Viewers"),
    color=alt.Color(
        "subscription_tier:N", scale=alt.Scale(scheme="purples"), legend=None
    ),
    tooltip=["subscription_tier", "count()"],
).properties(title="Viewer Distribution by Subscription Tier", width=400, height=300)
```

Out[117...



Average Watch Time Across Different Ages

Here, I decided to try and merge all the streaming data since I was successful before in order to speed up the data grouping process.

In [118...

```
# merge sessions with their viewer ids to see how many watch per session
# inner merge means only keys from both are kept (just in case i didnt clean good)
session_viewers = pd.merge(
    sessions_df_resolved,
    viewers_df_resolved,
    on="viewer_id",
    how="inner",
)

# combine it AGAIN with stream activity using stream_id to get full logistics
# we can't actually use this cause its way too big
full_stream_activity = pd.merge(
    session_viewers, streams_df_resolved, on="stream_id", how="inner"
)
```

In [119...

```
full_stream_activity.head(1)
```

Out[119...

	session_id	viewer_id	streamer_id_x	stream_id	started_at_x	ended_at_x	duration_mins	chat_message
0	SES000001	VWR00938	STR0055	SM00860	2024-01-18 19:16:00	2024-01-18 20:28:00	72	

1 rows × 24 columns

Then I focused on grouping and visualizing the age groups by duration

In [120...

```
# Group age groups by how long they watch streams on average
ages_avg_watchtime = (
    full_stream_activity.groupby("age_group")["duration_mins"]
    .mean()
    .reset_index()
)
```



```
)
```

In [121...

```
ages_avg_watchtime
```


Out[121...

	age_group	duration_mins
0	13-17	58.929220
1	18-24	59.151183
2	25-34	58.906152
3	35-44	60.195145
4	45+	59.016429

In [122...

```
alt.Chart(ages_avg_watchtime).mark_point(size=150, color = "purple").encode(
    y=alt.Y(
        "duration_mins:Q",
        title="Duration",
        axis=alt.Axis(format=",.0f"),
    ),
    x=alt.X("age_group:O", title="Age Group"),
    tooltip=[
        alt.Tooltip("age_group:N", title="Ages"),
        alt.Tooltip("duration_mins:Q", title="Average duration", format=",.0f"),
    ],
).properties(
    title="Average Watchtime by Age",
    width=500,
    height=300
)
```

Out[122...




2. Summary statistics for distribution of streams, peak viewership, and data overall

Total # of Streams by Stream Category

```
In [123... #unique means only count unique stream sessions, no duplicates (for instance more than one view)
total_stream_category = (
    full_stream_activity.groupby("category")["stream_id"]
    .nunique()
    .reset_index(name="stream_count")
)

alt.Chart(total_stream_category).mark_bar().encode(
    x=alt.X(
        "stream_count:Q",
        title="Total Number of Streams",
        axis=alt.Axis(format=",.0f"),
    ),
    y=alt.Y("category:N", title="Stream Category", sort="-x"),
    color=alt.Color(
        "category:N", scale=alt.Scale(scheme="purples"), legend=None
    ),
    tooltip=[
        alt.Tooltip("category:N", title="Category"),
        alt.Tooltip("stream_count:Q", title="Streams", format=",.0f"),
    ],
).properties(
    title="Total Streams by Stream Category",
    width=500,
    height=300,
)
```

Out[123...



In [124...

```
total_stream_category
```

Out[124...

	category	stream_count
0	Creative	191
1	FPS	208
2	IRL	127
3	Just Chatting	275
4	MOBA	170
5	Sports	178
6	unknown	50

Total Watch Time (Mins) of Viewers by Preferred Category

Used viewer_id to link duration_mins to viewer_id and then used that to display time for each preferred_category

- https://www.w3schools.com/python/pandas/ref_df_merge.asp

For reference, this was the code I wrote before creating `full_stream_activity` where I learned to merge and group data. I moved it to this section as it fit the instructions better.

In [125...

```
#create new dataframe column for viewer watchtime by linking viewer_id in viewers_df to duration_mins
viewer_watchtime = (
    sessions_df_resolved.groupby("viewer_id")["duration_mins"].sum().reset_index()
)

#left joining: preserve key order & keeps all keys from the Left DataFrame (which is viewers)
merged_viewers = pd.merge(viewers_df_resolved, viewer_watchtime, on="viewer_id", how="left")

# if viewer id is missing or doesnt match, fill with 0 mins (just in case)
merged_viewers["duration_mins"] = merged_viewers["duration_mins"].fillna(0)


# group by category
category_summary = (
    merged_viewers.groupby("preferred_category")["duration_mins"]
    .sum()
    .reset_index()
)
```

In [126...

```
alt.Chart(category_summary).mark_bar().encode(
    x=alt.X(
        "duration_mins:Q",
        title="Total Watch Time (Minutes)",
        axis=alt.Axis(format=",.0f"),
    ),
    y=alt.Y("preferred_category:N", title="Preferred Category", sort="-x"),
    color=alt.Color(
        "preferred_category:N", scale=alt.Scale(scheme="purples"), legend=None
    ),
    tooltip=[
        alt.Tooltip("preferred_category:N", title="Category"),
        alt.Tooltip("duration_mins:Q", title="Total Minutes", format=",.0f"),
    ],
).properties(
    title="Total Engagement (Watch Time) by Preferred Category",
```

```
width=500,  
height=300,)
```

Out [126...



Peak viewership for each month

- I used this stack overflow to extract the month names from pandas datetime object
 - <https://stackoverflow.com/questions/57033657/how-to-extract-month-name-and-year-from-date-column-of-dataframe>

In [127...


```
# Extract month name from our converted pandas date object  
sessions_df_resolved['month'] = sessions_df_resolved['started_at'].dt.strftime('%B')  
  
monthly_viewership = (sessions_df_resolved.groupby("month").agg(  
    total_sessions=("session_id", "count",), # distinct watch sessions  
    total_watch_hours=("duration_mins", lambda x: x.sum() / 60,), # Summing total watch hours  
) .reset_index())  
  
month_order = [  
    "January",  
    "February",  
    "March",  
    "April",  
    "May",  
    "June",  
    "July",  
    "August",  
    "September",  
    "October",  
    "November",  
    "December",  
]  
  
alt.Chart(monthly_viewership).mark_line(point=True, color="purple").encode(  
    x=alt.X(  
        "month:N",  
        title="Month (2024)",  
        sort=month_order, # Enforces correct calendar sequence  
        axis=alt.Axis(labelAngle=-45),
```

```

),
y=alt.Y(
    "total_watch_hours:Q",
    title="Total Viewership Volume (Hours)",
    axis=alt.Axis(format=",.0f"),
),
tooltip=[
    alt.Tooltip("month:N", title="Month"),
    alt.Tooltip("total_watch_hours:Q", title="Total Watch Hours", format=",.1f"),
    alt.Tooltip("total_sessions:Q", title="Total Sessions", format=","),
],
).properties(
    title="Total Platform Viewership Trend by Month",
    width=600,
    height=350)

```

Out[127...



3. Show category-specific information and compare streamers based on categories

"Just Chatting" Category by Language of Viewers

In [128...

```

just_chatting_streamers = streamers_df[streamers_df["category"] == "Just Chatting"]


alt.Chart(just_chatting_streamers).mark_bar().encode(
    x=alt.X(
        "sum(avg_concurrent_viewers):Q",
        title="Total Concurrent Viewers",
        axis=alt.Axis(format=",.0f"),
    ),
    y=alt.Y("language:N", title="Streamer Language", sort="-x"),
    color=alt.Color(
        "language:N", scale=alt.Scale(scheme="purples"), legend=None
    ),
    tooltip=[
        alt.Tooltip("language:N", title="Language"),

```

```

    alt.Tooltip(
        "sum(avg_concurrent_viewers):Q",
        title="Total Viewers",
        format=",.0f",
    ),
],
).properties(
    title="Just Chatting Category Viewers by Language Group",
    width=500,
    height=200,
).display()

```



Popularity Comparison Among MOBA Streamers


In [129...

```

moba_data = streamers_df[streamers_df["category"] == "MOBA"]

alt.Chart(moba_data).mark_point(size=120, filled=True).encode(
    x=alt.X(
        "total_followers:Q",
        title="Total Followers",
        axis=alt.Axis(format=",.0s"),
    ),
    y=alt.Y(
        "avg_concurrent_viewers:Q",
        title="Avg Concurrent Viewers",
        axis=alt.Axis(format=",.0s"),
    ),
    color=alt.Color("streamer_name:N"),
    tooltip=[
        alt.Tooltip("streamer_name:N", title="Streamer"),
        alt.Tooltip("total_followers:Q", title="Followers", format=","),
        alt.Tooltip(
            "avg_concurrent_viewers:Q", title="Avg Viewers", format=","
        ),
    ],
).properties(
    title="MOBA Streamer Popularity Comparison",
    width=550,
    height=350,
).display()

```



Popularity of 'FPS' vs 'Just Chatting' streamers


Using average concurrent viewers, we can calculate the mean/median/max of each stream category and then boxplot them to see where the extremes can reach vs. average

- Altair automatically calculates this for us in the tooltip

In [130...

```
fps_vs_jc = streamers_df[streamers_df["category"].isin(["FPS", "Just Chatting"])]

alt.Chart(fps_vs_jc).mark_boxplot(extent="min-max", size=40).encode(
    x=alt.X("category:N", title="Stream Category"),
    y=alt.Y(
        "avg_concurrent_viewers:Q",
        title="Average Concurrent Viewers",
        axis=alt.Axis(format=",.0f"),
    ),
    color=alt.Color(
        "category:N",
        scale=alt.Scale(domain=["FPS", "Just Chatting"], scheme="purples"),
        legend=None,
    )
).properties(
    title="FPS vs. Just Chatting Streamers",
    width=400,
    height=400,
).display()
```



Task 3 Summary report

1. User Characteristics & Age Demographics

- **Demographics by Country:** display the distribution of viewer age groups across different countries.
- **Subscription Tiers:** Isolated viewer counts by grouping categorical tiers using `sort=` to distinguish the free accounts through paid tiers (`free` -> `tier1` -> `tier2` -> `tier3`).
- **Average Watch Time by Age:** shared relational columns using an inner `.merge()` on `viewer_id` and `stream_id` . Grouped categories with `.groupby()` and calculated mean with `.mean()`

2. Stream Distribution & Overall Platform Metrics

- **Stream Category Sizing:** using `.groupby()` on categories combined with `.nunique()` to count unique instances of `stream_id`
- **Preferred Category Engagement:** Got total viewer engagement by calculating `.sum()` of total watch time per viewer. Merged them back before mapping the totals to a horizontal bar chart.
- **Chronological Trends:** Used `.dt.strftime('%B')` to isolate month names. I divided overall minutes by 60 to convert them into a hours scale using `.agg()`

3. Category-Specific Comparisons & Streamer Analytics

- **Just Chatting Languages:** Filtered `streamers_df.csv` to content creators under the `Just Chatting`

genre. Divided viewership metrics by language.

- **MOBA Peer Performance:** Extracted streamer popularity metrics within the `MOBA` category. Used a plot map to easily highlight top-performing creator outliers.
- **FPS vs. Just Chatting Benchmarks:** Used the method `.isin()` to filter down and compare the dataset's two largest genres.

Task 4: Dealing with Outliers and Missing Values

I had some help with this section with this article by Hasan Ersan

- <https://hersanyagci.medium.com/detecting-and-handling-outliers-with-pandas-7adbcd5cad8>

Missing Values / Outliers in the original data

Analysis (see **Task 1**):

- `sessions.csv` = `streamer_id` has 100 null entries
- `streamers.csv` = `language` has 8 null entries
- `streams.csv` = `category` has 50 null entries
- `viewers.csv` = `country` has 80 null entries

I already dealt with these problems in **Task 2** so I'll just reiterate what I've done to explain my two strategies

`viewers.csv`

- Missing values for 'country' (80)

Strategy 1: Resolved by adding new country label 'unknown' to stay specific but not drop potentially important entries.

In [131...

```
viewers_df_resolved = viewers_df.copy()

viewers_df_resolved['country'] = viewers_df_resolved['country'].fillna('unknown')

viewers_df_resolved.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   viewer_id             1000 non-null  str
1   age_group             1000 non-null  str
2   country               1000 non-null  str
3   account_age_days      1000 non-null  int64
4   subscription_tier      1000 non-null  str
5   preferred_category     1000 non-null  str
dtypes: int64(1), str(5)
memory usage: 47.0 KB
```

streams.csv

- 'peak_viewers' has 80 null entries

Strategy 2: Resolved by replacing it with the median instead of mean due to outliers being present in the form of viral streams.

```
In [132... # Please see code in Task 2 I don't want to accidentally run it twice
```

Missing Values / Outliers in the derived data

First I use this function we defined in Lab02 to find outliers in the data using the +/- std

```
In [133... def detect_outliers(df, column):
    mean = df[column].mean() # the mean = mean of that specific data column
    std = df[column].std() # standard deviation
    lower_bound = mean - 2 * std # mean -2 std
    higher_bound = mean + 2 * std # mean +2 std
    return df[(df[column] < lower_bound) | (df[column] > higher_bound)]
```

streams_df_resolved

- there are around 61 extremely high performing streams that don't really reflect the average peak twitch viewership count on a session

Strategy 1: Resolved by removing the outliers entirely

```
In [134... # First analyze the streams data using the function


streams_outliers = detect_outliers(streams_df_resolved, "peak_viewers")
streams_outliers.info()
print()
print(streams_outliers['peak_viewers'])
```

```
<class 'pandas.DataFrame'>
Index: 64 entries, 2 to 979
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   stream_id             64 non-null    str
1   streamer_id           64 non-null    str
2   started_at            64 non-null    datetime64[us]
3   ended_at              64 non-null    datetime64[us]
4   stream_duration_hrs   64 non-null    float64
5   category              64 non-null    str
6   peak_viewers          64 non-null    float64
7   title_has_hype_word   64 non-null    bool
8   was_raided            64 non-null    bool
dtypes: bool(2), datetime64[us](2), float64(2), str(3)
memory usage: 4.1 KB
```

```
2      212127.0
3      186745.0
8      161545.0
11     148097.0
132    165095.0
...
969    296232.0
970    299413.0
973    210085.0
976    338648.0
979    330664.0
Name: peak_viewers, Length: 64, dtype: float64
```

```
In [135... # Visualize it with boxplot just to see if its accurate

streams_df_resolved.boxplot(column='peak_viewers', figsize=(6,4))
plt.show()
```



This is where I remove these extremely well-performing streams using their indexes

```
In [136... streams_outliers_removed = streams_df_resolved.drop(streams_outliers.index)
print(streams_outliers_removed)
```

	stream_id	streamer_id		started_at		ended_at	\
0	SM00001	STR0001	2024-02-14	18:00:00	2024-02-14	21:00:00	
1	SM00002	STR0001	2024-01-28	18:00:00	2024-01-28	21:00:00	
4	SM00005	STR0001	2024-01-24	18:00:00	2024-01-24	21:00:00	
5	SM00006	STR0001	2024-01-09	14:00:00	2024-01-09	17:00:00	
6	SM00007	STR0001	2024-03-05	12:00:00	2024-03-05	16:00:00	
...	...	...		...		...	
1194	SM01195	STR0076	2024-01-31	19:00:00	2024-01-31	23:00:00	
1195	SM01196	STR0076	2024-03-20	19:00:00	2024-03-20	21:30:00	
1196	SM01197	STR0076	2024-01-01	14:00:00	2024-01-01	21:00:00	
1197	SM01198	STR0076	2024-01-02	20:00:00	2024-01-02	23:00:00	
1198	SM01199	STR0076	2024-02-19	08:00:00	2024-02-19	10:00:00	

	stream_duration_hrs	category	peak_viewers	title_has_hype_word	\
0	3.0	Just Chatting	108989.0	False	
1	3.0	Just Chatting	51635.0	True	
4	3.0	Just Chatting	87723.0	False	
5	3.0	Just Chatting	74062.0	True	
6	48.0	Just Chatting	20224.0	True	
...	...	...	...	...	
1194	4.0	FPS	1421.0	False	
1195	2.5	FPS	2032.0	True	
1196	7.0	FPS	1967.0	False	
1197	3.0	FPS	2041.0	False	
1198	2.0	FPS	771.0	False	

	was_raid
0	True
1	False
4	True
5	False
6	False
...	...
1194	False
1195	False
1196	False
1197	False
1198	False

[1135 rows x 9 columns]

```
In [137... print(f"Original dataset size: {len(streams_df_resolved)}")
print(f"After removing outliers: {len(streams_outliers_removed)}")
```

Original dataset size: 1199
After removing outliers: 1135

```
In [138... streams_outliers_new = detect_outliers(streams_outliers_removed, "peak_viewers")
streams_outliers_new.info()
print()
print(streams_outliers_new['peak_viewers'])
```

```
<class 'pandas.DataFrame'>
Index: 76 entries, 0 to 1131
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   stream_id              76 non-null    str
1   streamer_id            76 non-null    str
2   started_at             76 non-null    datetime64[us]
3   ended_at               76 non-null    datetime64[us]
4   stream_duration_hrs    76 non-null    float64
5   category               76 non-null    str
6   peak_viewers           76 non-null    float64
7   title_has_hype_word    76 non-null    bool
8   was_raided             76 non-null    bool
dtypes: bool(2), datetime64[us](2), float64(2), str(3)
memory usage: 4.9 KB
```

```
0      108989.0
4       87723.0
7      118974.0
9      131075.0
57     107512.0
...
1114     92752.0
1118     102761.0
1123     107908.0
1128      91237.0
1131      87681.0
```

```
Name: peak_viewers, Length: 76, dtype: float64
```

By removing the outliers, the standard deviation gets reduced but also unfortunately permanently deletes those viral streams with high amount of peak viewers. Now peak viewers never exceed 200000.

streamers_df

- there are around 5 streamers with extremely high total follower counts > 11100000 followers

Strategy 2: Normalize the follower count using log transformation to understand ratios/magnitude better

In [139...

```
# First analyze the streamers data using the outlier function
streamer_outliers = detect_outliers(streamers_df, "total_followers")
streamer_outliers.info()
print()
print(streamer_outliers['total_followers'])
```

```
<class 'pandas.DataFrame'>
Index: 5 entries, 0 to 61
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   streamer_id            5 non-null     str
1   streamer_name          5 non-null     str
2   category               5 non-null     str
3   language               5 non-null     str
4   partner_status         5 non-null     bool
5   total_followers        5 non-null     int64
6   avg_concurrent_viewers 5 non-null     int64
dtypes: bool(1), int64(2), str(4)
memory usage: 285.0 bytes
```

```
0      11500000
8      14500000
18     11100000
60     12300000
61     18700000
Name: total_followers, dtype: int64
```

```
In [140... streamers_df.boxplot(column='total_followers', figsize=(6,4))
plt.show()
```




Apply the log transformations here and then check the boxplot again

```
In [141... streamers_log = np.log1p(streamers_df["total_followers"])
streamers_log.head(10)
```

```
Out[141... 0    16.257858
1    10.148471
2    14.771022
3    15.640060
4    14.557448
5    13.764218
6    10.717325
7    10.757158
8    16.489659
9    16.087637
Name: total_followers, dtype: float64
```

```
In [142... #create new column for log transformations in dataset
streamers_df["log_total_followers"] = np.log1p(streamers_df["total_followers"])
streamers_df.boxplot(column="log_total_followers", figsize=(6, 4))
```

```
Out[142... <Axes: >
```



```
In [143... # Check new outliers in the logged data
streamer_log_outliers = detect_outliers(streamers_df, "log_total_followers")
streamer_log_outliers.info()
print()
print(streamer_log_outliers['log_total_followers'])
```

```
<class 'pandas.DataFrame'>
Index: 5 entries, 1 to 24
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   streamer_id            5 non-null     str
1   streamer_name          5 non-null     str
2   category               5 non-null     str
3   language               5 non-null     str
4   partner_status         5 non-null     bool
5   total_followers        5 non-null     int64
6   avg_concurrent_viewers 5 non-null     int64
7   log_total_followers    5 non-null     float64
dtypes: bool(1), float64(1), int64(2), str(4)
memory usage: 325.0 bytes
```

```
1    10.148471
6    10.717325
7    10.757158
12    9.094144
24    10.495045
```

```
Name: log_total_followers, dtype: float64
```

Instead of removing streamers with large followings, we just rescale the distribution of total_followers. This reflects magnitude/scale rather than literal count.

Task 4 Summary report

Please see Task 2 to see a summary of the entire cleaning and other imputation processes

Handling Missing Values in the Original Data

- **Viewer Country Fix (Strategy 1 - Placeholder):** The original user profile dataset had 80 blanks in the country column. Dropping these records risked losing valuable engagement metrics from other columns. This was resolved by applying `.fillna('unknown')` to insert a clean placeholder, preserving all 1,000 original profile entries.
- **Stream Peak Viewers Fix (Strategy 2 - Median Imputation):** The broadcast history dataset contained 80 missing entries for concurrent viewership. Because the platform contains massive viral stream spikes, using a standard average would have pulled values upward unfairly. This was resolved by substituting the blanks with the statistical median value of the category instead.

Managing Outliers in Derived Data

Strategy 1: Removing Extreme Outliers (Peak Viewership)

- **Process:** Running a standard deviation-based helper function revealed 64 hyper-viral streaming sessions where peak metrics sat multiple standard deviations above normal platform benchmarks. A boxplot visual confirmed these extreme values stretched heavily toward the top end. The records were extracted and physically discarded from the final working dataset using `.drop()`.
- **Result:** This dropped the total broadcast row count from 1,199 down to 1,135 streams. While it drastically lowered the data's standard deviation and capped the maximum viewer scale at a predictable 200,000 baseline, it permanently erased the historical footprint of major viral events from subsequent calculations.

Strategy 2: Log Transformation (Total Follower Counts)

- **Process:** The initial data scan isolated 5 massive platform celebrities with exceptional follower milestones exceeding 11 million users, creating a heavily skewed boxplot. Instead of deleting these highly important creators, a mathematical log transformation was introduced using `np.log1p()`. This new calculation was stored inside a distinct DataFrame column.
- **Result:** Rather than losing data points, this process squeezed the massive, multi-million follower gaps into a symmetrical range between roughly 9 and 16. Re-running the boxplot tool over this transformed scale confirmed that the distribution became normalized, switching the focal perspective to percentage-based magnitude and scale rather than raw numbers."""

Task 5: Correlation and Behavioral Patterns

1. Compute pairwise correlations between user features and the number of check-ins

I'm not too sure what this means. I assume it means correlate user features like tier and the activity of users like watching streams/sessions, giving bits, and sending chat messages.

In [144...

```
# count number of sessions and sum up the activity of the user in that id
user_activity = (sessions_df_resolved.groupby("viewer_id").agg(
    sessions=("session_id", "count"),
    total_watch_time=("duration_mins", "sum"),
    chat_messages=("chat_messages_sent", "sum"),
    bits_cheered=("bits_cheered", "sum"),
).reset_index()
)

# merge with the user 'features' like tier, country, & age group
# inner only keeps the entry if the viewer id is in both viewers.csv and our user_activity
features_df = pd.merge(
    viewers_df_resolved, user_activity, on="viewer_id", how="inner"
)

#preview & double check theres no null columns
features_df.head()
features_df.info()

#select these columns only to put into a new correlation dataframe
correlation_cols = [
    "account_age_days",
    "sessions",
    "total_watch_time",
    "chat_messages",
    "bits_cheered",
]

correlation_data = features_df[correlation_cols]
```

```
<class 'pandas.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   viewer_id             1000 non-null   str
1   age_group             1000 non-null   str
2   country               1000 non-null   str
3   account_age_days      1000 non-null   int64
4   subscription_tier      1000 non-null   str
5   preferred_category     1000 non-null   str
6   sessions              1000 non-null   int64
7   total_watch_time       1000 non-null   int64
8   chat_messages         1000 non-null   int64
9   bits_cheered          1000 non-null   int64
dtypes: int64(5), str(5)
memory usage: 78.3 KB
```

Compute basic correlations just to see how the data affects eachother in easy to read raw format.

In [145...

```
corr_matrix = features_df[correlation_cols].corr()

print(corr_matrix.round(2))
```

	account_age_days	sessions	total_watch_time	chat_messages	\
account_age_days	1.00	-0.02	-0.02	0.02	
sessions	-0.02	1.00	0.98	0.59	
total_watch_time	-0.02	0.98	1.00	0.59	
chat_messages	0.02	0.59	0.59	1.00	
bits_cheered	-0.03	0.80	0.79	0.45	

	bits_cheered
account_age_days	-0.03
sessions	0.80
total_watch_time	0.79
chat_messages	0.45
bits_cheered	1.00

2. Visualize the results using a correlation matrix (use a coolwarm color scheme and proper scaling)


Pearson Correlation Matrix

I only realised the default correlation matrix is Pearson after I finished this code, hence why I made `pearson_corr` instead of using the correlation data directly. That means the raw text matrix from before is also the Pearson method and a "default" correlation pandas method doesn't actually exist (its just Pearson)

In [146...

```
#pearson corelation method
pearson_corr = correlation_data.corr(method="pearson")

plt.figure(figsize=(8, 6))
sns.heatmap(pearson_corr, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Pearson Correlation Matrix", fontsize=14)
plt.tight_layout()
plt.show()
```

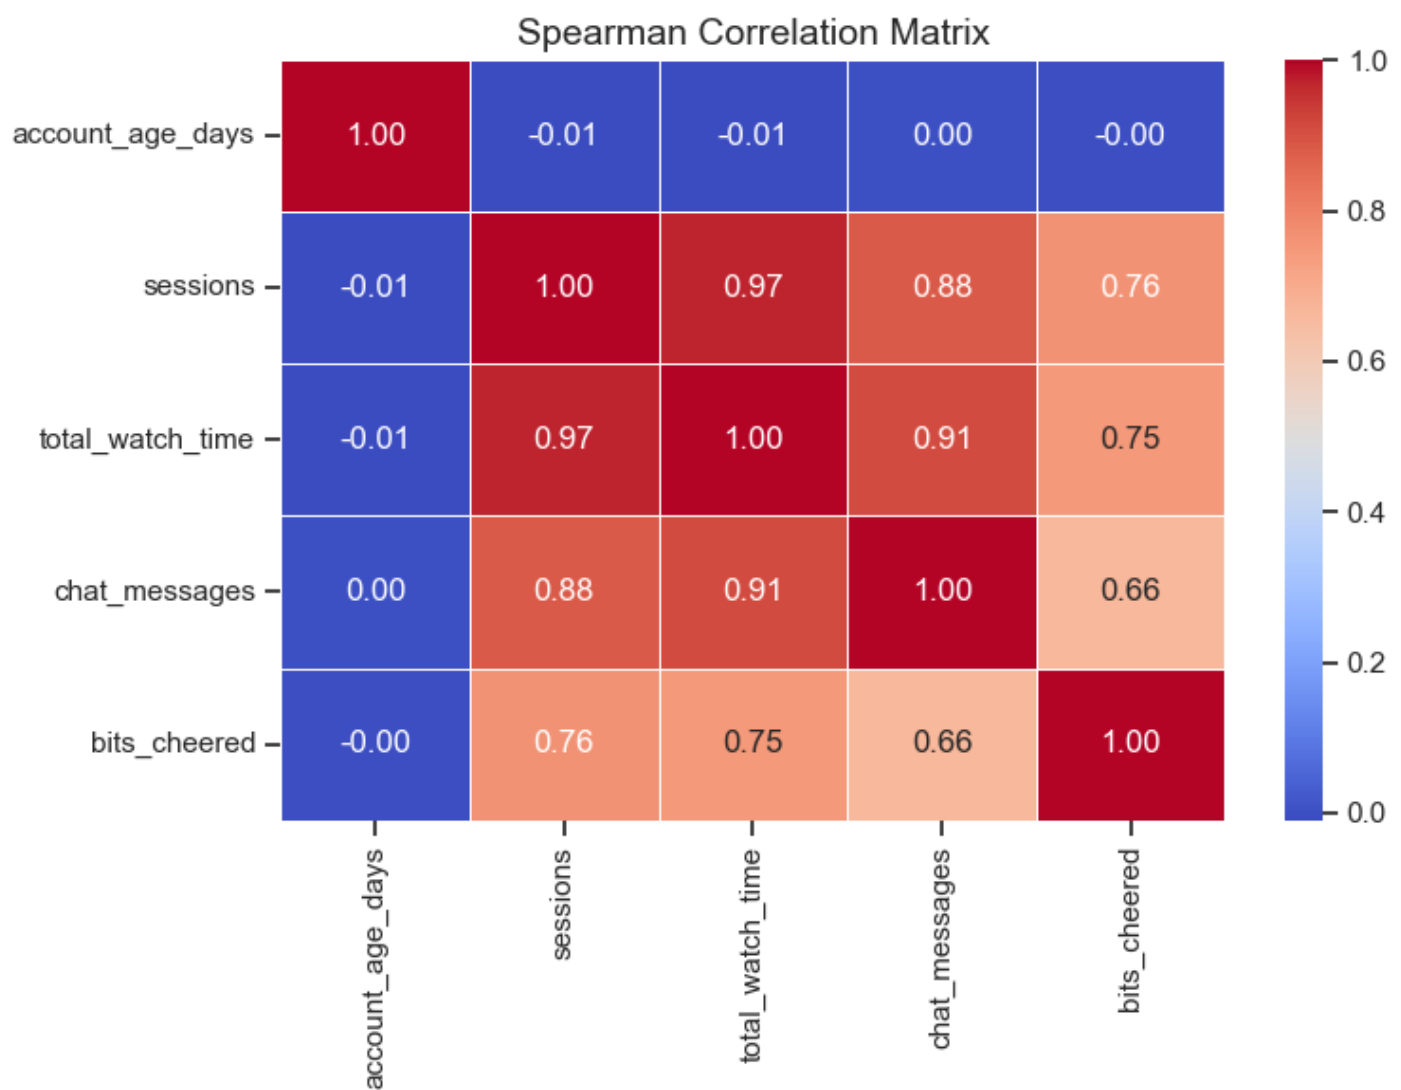


Spearman Correlation Matrix

In [147...

```
spearman_corr = correlation_data.corr(method="spearman")

plt.figure(figsize=(8, 6))
sns.heatmap(spearman_corr, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Spearman Correlation Matrix", fontsize=14)
plt.tight_layout()
plt.show()
```



3. Try Spearman rank correlation and explain when it is more appropriate than Pearson

Pearson measures **linear relationships** better, so if the data goes up in a straight diagonal line (positive correlation or negative correlation only) it is the ideal choice.

Otherwise the Spearman method measures **monotonic fit** and is stronger/less sensitive against outliers. So if the relationship is non-linear, where values go up and down together regardless of correlation pattern, then Spearman is better.

For this dataset, **Spearman is better for the correlation matrix**. If some extreme fans cheer an extreme amount of bits, Spearman will rank (ex. this is #1 highest bits) it instead of comparing the stats as its literal number. It also works better for ordinal data, like tiers 1-3.

Pearson would be really sensitive to these extreme numbers as it counts the numbers literally without fully taking in the context of the correlation, it just correlates them.

Task 5 Summary report

1. User Characteristics & Age Demographics

- **Demographics by Country:** display the distribution of viewer age groups across different countries.

- **Subscription Tiers:** Isolated viewer counts by grouping categorical tiers using `sort=` to distinguish the free accounts through paid tiers (`free` -> `tier1` -> `tier2` -> `tier3`).
- **Average Watch Time by Age:** shared relational columns using an inner `.merge()` on `viewer_id` and `stream_id` . Grouped categories with `.groupby()` and calculated mean with `.mean()`

2. Stream Distribution & Overall Platform Metrics

- **Stream Category Sizing:** using `.groupby()` on categories combined with `.nunique()` to count unique instances of `stream_id`
- **Preferred Category Engagement:** Got total viewer engagement by calculating `.sum()` of total watch time per viewer. Merged them back before mapping the totals to a horizontal bar chart.
- **Chronological Trends:** Used `.dt.strftime('%B')` to isolate month names. I divided overall minutes by 60 to convert them into a hours scale using `.agg()`

Task 6: Compare Behavior Across User Segments

Monthly watchtime by age-group

In [148...

```
full_stream_activity.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 22679 entries, 0 to 22678
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   session_id            22679 non-null  str
1   viewer_id             22679 non-null  str
2   streamer_id_x         22679 non-null  str
3   stream_id             22679 non-null  str
4   started_at_x          22679 non-null  datetime64[us]
5   ended_at_x            22679 non-null  datetime64[us]
6   duration_mins         22679 non-null  int64
7   chat_messages_sent    22679 non-null  int64
8   bits_cheered          22679 non-null  int64
9   followed_during       22679 non-null  bool
10  subscribed_during      22679 non-null  bool
11  age_group              22679 non-null  str
12  country                22679 non-null  str
13  account_age_days      22679 non-null  int64
14  subscription_tier      22679 non-null  str
15  preferred_category     22679 non-null  str
16  streamer_id_y         22679 non-null  str
17  started_at_y          22679 non-null  datetime64[us]
18  ended_at_y            22679 non-null  datetime64[us]
19  stream_duration_hrs   22679 non-null  float64
20  category              22679 non-null  str
21  peak_viewers          22679 non-null  float64
22  title_has_hype_word    22679 non-null  bool
23  was_raided            22679 non-null  bool
dtypes: bool(4), datetime64[us](4), float64(2), int64(4), str(10)
memory usage: 3.5 MB
```

Article that helped me understand how to properly group data together:

- <https://www.geeksforgeeks.org/pandas/python-pandas-dataframe-groupby/>

In [166...

```
#Merge sessions_df with viewers_df to get full behaviour stats
viewer_sessions_df = pd.merge(viewers_df_resolved, sessions_df_resolved, on="viewer_id", how="inner")

# Group by age group & month then count individual sessions + add watchtime
aggregated_viewer_sessions_month = viewer_sessions_df.groupby(['month', 'age_group']).agg(
    total_sessions=("session_id", "count"), # distinct watch sessions
    total_watch_hours=("duration_mins", lambda x: x.sum() / 60), # Summing total watch hours
).reset_index()

print(aggregated_viewer_sessions_month)
```

	month	age_group	total_sessions	total_watch_hours
0	February	13-17	615	647.350000
1	February	18-24	2519	2383.716667
2	February	25-34	2028	1900.316667
3	February	35-44	973	1015.716667
4	February	45+	827	803.000000
5	January	13-17	760	763.883333
6	January	18-24	2931	2832.150000
7	January	25-34	2499	2448.383333
8	January	35-44	1142	1102.400000
9	January	45+	1036	998.800000
10	March	13-17	716	642.450000
11	March	18-24	2494	2615.750000
12	March	25-34	2186	2241.916667
13	March	35-44	1016	1023.066667
14	March	45+	937	952.300000

In [167...

```
month_order = [
    "January",
    "February",
    "March",
    "April",
    "May",
    "June",
    "July",
    "August",
    "September",
    "October",
    "November",
    "December",
]


alt.Chart(aggregated_viewer_sessions_month).mark_circle(size=100).encode(
    x=alt.X(
        "month:N",
        title="Month (2024)",
        sort=month_order,
        axis=alt.Axis(labelAngle=-45),
    ),
    y=alt.Y(
        "total_watch_hours:Q",
        title="Total Viewership (Hours)",
        axis=alt.Axis(format=",.0f"),
    ),
    color='age_group:N',
    tooltip=[
        alt.Tooltip("age_group:N", title="Age Group"),
        alt.Tooltip("month:N", title="Month"),
    ]
)
```

```

alt.Tooltip("total_watch_hours:Q", title="Total Watch Hours", format=",.1f"),
alt.Tooltip("total_sessions:Q", title="Total Sessions", format=","),
],
).properties(
    title="Monthly watchtime by Age-Group",
    width=600,
    height=350)

```

Out[167...



Monthly Watch Time Between Free & Tier Accounts

In [168...

```

# Seperate subscription tier into sub group of paid / free
viewer_sessions_df["sub_group"] = viewer_sessions_df["subscription_tier"].apply(lambda x: "Free"

# aggregate session duration sum by month, user group, and unique viewer
sub_monthly_watch = viewer_sessions_df.groupby(["month", "sub_group", "viewer_id"])["duration_mir

# turn minutes into hours
sub_monthly_watch["watch_hours"] = sub_monthly_watch["duration_mins"] / 60

sub_monthly_watch

```


Out[168...

	month	sub_group	viewer_id	duration_mins	watch_hours
0	February	Free	VWR00002	44	0.733333
1	February	Free	VWR00003	272	4.533333
2	February	Free	VWR00005	62	1.033333
3	February	Free	VWR00010	17	0.283333
4	February	Free	VWR00013	117	1.950000
...	...	...	...	...	...
2579	March	Paid (Tier 1-3)	VWR00989	320	5.333333
2580	March	Paid (Tier 1-3)	VWR00991	1137	18.950000
2581	March	Paid (Tier 1-3)	VWR00992	1917	31.950000
2582	March	Paid (Tier 1-3)	VWR00996	1672	27.866667
2583	March	Paid (Tier 1-3)	VWR00998	2379	39.650000

2584 rows × 5 columns

In [152...

```
alt.Chart(sub_monthly_watch).mark_boxplot(extent="min-max", size=25).encode(
    x=alt.X("sub_group:N", title="Subscription Status"),
    y=alt.Y(
        "watch_hours:Q",
        title="Monthly Watch Time per User (Hours)",
        axis=alt.Axis(format=",.1f"),
    ),
    color=alt.Color(
        "sub_group:N",
        scale=alt.Scale(domain=["Free", "Paid (Tier 1-3)"], scheme="set2"),
        legend=None,
    ),
    column=alt.Column(
        "month:N", title="Timeline (2024)", sort=month_order, spacing=15
    ),
).properties(
    title="Monthly Watch Time Between Free & Tier Accounts",
    width=200,
    height=300,
)
```

Identify any outliers in group behavior and investigate them

I decided to analyze any outlier behaviour between the watch session duration of free users vs. paid ones.

These ended up being anyone who watched a stream for longer than >27 hrs.

In [153...

```
# run the std outliers function on the dataset
sub_monthly_outliers = detect_outliers(sub_monthly_watch, "watch_hours")
sub_monthly_outliers.info()
print()
print(sub_monthly_outliers['watch_hours'])
print()

#check how much of the original dataset is affected
print(f"Original dataset size: {len(sub_monthly_watch)}")
print(f"Outliers size: {len(sub_monthly_outliers)}")
```

```
<class 'pandas.DataFrame'>
Index: 155 entries, 420 to 2583
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   month            155 non-null    str
1   sub_group        155 non-null    str
2   viewer_id        155 non-null    str
3   duration_mins    155 non-null    int64
4   watch_hours      155 non-null    float64
dtypes: float64(1), int64(1), str(3)
memory usage: 7.3 KB
```

```
420    34.850000
426    32.216667
427    37.983333
428    27.150000
440    31.083333
```

```
...
2571   28.866667
2576   42.583333
2581   31.950000
2582   27.866667
2583   39.650000
```

Name: watch_hours, Length: 155, dtype: float64

Original dataset size: 2584

Outliers size: 155

In [154...

```
# check the minimum hours for the outliers
print()
print(sub_monthly_outliers.sort_values(by="watch_hours", ascending=True).head(10))
```

	month	sub_group	viewer_id	duration_mins	watch_hours
1297	January	Free	VWR00889	1625	27.083333
2176	March	Free	VWR00896	1625	27.083333
2191	March	Free	VWR00923	1626	27.100000
428	February	Free	VWR00865	1629	27.150000
1694	January	Paid (Tier 1-3)	VWR00937	1631	27.183333
1667	January	Paid (Tier 1-3)	VWR00866	1643	27.383333
2216	March	Free	VWR00967	1649	27.483333
1707	January	Paid (Tier 1-3)	VWR00964	1657	27.616667
2217	March	Free	VWR00968	1660	27.666667
2167	March	Free	VWR00881	1662	27.700000

A majority of the outliers seem to be free accounts. This means we could potentially investigate why they choose to stay a free account despite their high watchtime.

Task 6 Summary report

Age-Based User Behaviour & Comparison

- **User Group Merging:** Combined user profiles and session tracking logs using an inner `.merge()` on the shared tracking ID column as well as `.agg()` to calculate combined watch time.
- **Timeline Visualization:** Plotted the aggregated on a dot plot (`mark_circle`), mapping cumulative viewing hours based on both `age_group` and `subscription_tier` (subgrouped into free + paid). Young adults aged 18–24 consistently generated the highest total viewership volumes across the tracked months.

Behavioral Outliers

- **Subgroup Tiers:** Separated users into two clear consumer groups, "Free" accounts and "Paid (Tier 1-3)" subscribers.
- **Distribution Mapping:** Plotted these distributions as sequential boxplots using `mark_boxplot` to evaluate visual spread over 3 months.
- **Outlier Investigation:** Ran a standard deviation outlier detection helper function on the user metrics, isolating 155 entries. Surprisingly, a major portion of these super-consumers were using entirely free accounts meaning they can benefit from targeted subscription upgrade promotions.

Task 7: Your Own EDA Questions

Hypthosis 1: Does title hype-word usage correlate to higher peak viewer count?

Group `streams_df_resolved` on whether the title has a hype word or not and then aggregate the peak viewers using built in math functions.


```
In [155... # group whether title has a hype word or not and then aggregate the peak viewers using built in m
hype_title_summary = (streams_df_resolved.groupby("title_has_hype_word")["peak_viewers"]
    .agg(["count", "mean", "median", "max"]).reset_index())

hype_title_summary
```

```
Out[155... 
```

	title_has_hype_word	count	mean	median	max
0	False	853	37680.069168	20224.0	370165.0
1	True	346	35893.592486	20370.0	272708.0

```
In [156... alt.Chart(streams_df_resolved).mark_boxplot(
    extent="min-max", size=50
).encode(
    x=alt.X(
        "title_has_hype_word:N",
        title="Title Has Hype Word?",
        axis=alt.Axis(labelAngle=0),
    ),
    y=alt.Y(
        "peak_viewers:Q",
        title="Peak Viewers",
        axis=alt.Axis(format=",.0f"),
    ),
    color=alt.Color(
        "title_has_hype_word:N",
        scale=alt.Scale(scheme="purples"),
        legend=None,
    ),
).properties(
    title="Impact of Title Hype Words on Peak Viewership", width=400, height=350)
```



Conclusion

Surprisingly, clickbait "hype word" in the title does not drive extra traffic. The max viewership of streams without the words, `false`, actually exceeds those that are `true`. The average viewership is slightly higher as well. Therefore hypewords don't correlate with getting more viewers and seems to have the opposite effect.

- Though the Q1 and the median are higher when `true` the Q3 isn't. This might suggest outliers are skewing the data.

Hypthosis 2: Are morning or evening streams more popular with viewers and what category do they belong to?

First I have to do the same thing I did last time when I extracted the month but with the hour of the day this time. Since we use a 24 hour timeframe, we can just segment the pandas datetime object into different sections of the day.

- After that we group them by category and aggregate the viewership to them

In [157...

```
# create new stream hours column
streams_df_resolved["stream_hour"] = streams_df_resolved["started_at"].dt.hour

# function to automatically sub-group them without doing it manually
# since it's 24 hour time it's easy to seperate them into groups
# Morning = 5AM-12PM
# Evening = 5-11PM

def assign_time_of_day(hour):
    if 5 <= hour < 12:
        return "Morning"
    elif 17 <= hour <= 23:
        return "Evening"
```

```

    else:
        return "Other"

# use the function to assign time of day to every stream in new sub group column
streams_df_resolved["time_of_day"] = streams_df_resolved["stream_hour"].apply(assign_time_of_day)

# Group time of day and category together, then aggregate peak viewers to them
time_category_summary = (streams_df_resolved.groupby(["time_of_day", "category"])["peak_viewers"]
    .agg(["count", "mean", "median"])
    .reset_index())

# #Preview
time_category_summary.head(2)

```

Out[157...


	time_of_day	category	count	mean	median
0	Evening	Creative	76	15663.276316	13484.0
1	Evening	FPS	103	42797.097087	32405.0

In [169...

```

alt.Chart(time_category_summary).mark_circle(size=120).encode(
    x=alt.X(
        "time_of_day:N",
        axis=alt.Axis(labelAngle=-45),
    ),
    y=alt.Y(
        "mean:Q",
        title="Average Peak Viewers",
        axis=alt.Axis(format=",.0f"),
    ),
    color=alt.Color(
        "category:N",
        scale=alt.Scale(scheme="dark2"),
    ),
    tooltip=[
        alt.Tooltip("time_of_day:N", title="Time"),
        alt.Tooltip("category:N", title="Category"),
        alt.Tooltip("mean:Q", title="Avg Peak Viewers", format=",.0f"),
        alt.Tooltip("count:Q", title="Total Streams"),
    ],
).properties(
    title="Popularity of Stream Categories Morning vs. Evening",
    width=550,
    height=350)

```



Conclusion

Morning streams perform better overall for almost every category except **Creative**. **Just Chatting** reaches the highest peak viewership of any category during the morning. **IRL** viewership drops significantly past 11pm.

- Morning viewership seems to be the most popular and common.
- Meanwhile variety **unknown** streams surpass **Just Chatting** after 11pm.

Briefly discuss what additional data would help you answer the questions more effectively

To better evaluate Hype Titles: Right now, we only see the final peak viewership, but we don't know how many users saw the title in their dashboard feed and explicitly chose to ignore it vs. click on it. Maybe some sort of algorithm data would help us understand how many viewers were exposed to these streams.

To better evaluate Stream Timing: Because Twitch is a global platform, a stream that begins in the morning for a European creator is broadcasting live to midnight in America. Matching the viewers' timezone to the Streamer's timezone would add a lot of context

- Though I don't know how this could be graphed or visualized.

Task 7 Summary report

1. Hypothesis 1: Title Hype Words vs. Peak Viewers

- **Data Processing:** Grouped the broadcast data based on whether a stream title included a clickbait hype word or not using `.groupby()`. Calculated the total count, mean, median, and maximum peak viewers using `.agg()`.
- **Findings:** Surprisingly, using hype words in the title does not bring in more viewers. Streams without

hype words actually had a higher maximum viewership and a slightly higher average viewer count. While streams with hype words had a slightly higher median, the upper boundaries suggest that standard outliers might be twisting the raw metrics.

2. Hypothesis 2: Stream Timing and Category Popularity

- **Data Processing:** Extracted the raw hour from the stream start times using `.dt.hour`. Built a custom Python function to automatically sort these hours into "Morning", "Evening", or "Other" buckets, and applied it to the dataset using `.apply()`. Grouped the data by both time of day and stream genre to find the average viewer counts.
- **Findings:** The chart showed that morning streams actually do better overall for almost every single category except for Creative streams. `Just Chatting` brings in the biggest morning crowds, `IRL` drops off heavily late at night, and unclassified variety streams take over the top spot after 11 PM.

3. Future Data Improvements

- **Hype Word Tracking:** To really understand if clickbait titles work, it would help to have algorithm data. We only see the final viewer peaks but we do not know how many users were exposed to the stream.
- **Time Zone Adjustments:** Having the physical location or time zone of the viewers would make the stream timing data much more accurate.

Task 8: EDA Presentation (Insight Report)

Summarize your findings in a polished format:

- Create 3–4 slides or a single-page report highlighting your most important insights
- Include at least 2 charts and interpret them using clear, plain language
- Apply Tufte's principles: minimize chartjunk, maximize data-to-ink ratio, and ensure clear labeling

Please see .pdf for the single-page report

Edward Tufte Design Principles

In [159...

```
plt.figure(figsize=(6, 4))

# convert the dataframe into matrix format for easy manipulation/display
# and define labels
labels = ["Account Age", "Sessions", "Watch Time", "Chat Sent", "Bits Cheered"]
matrix_data = np.array(
    [
        [1.00, -0.01, -0.01, 0.00, -0.00],
        [-0.01, 1.00, 0.97, 0.88, 0.76],
        [-0.01, 0.97, 1.00, 0.91, 0.75],
        [0.00, 0.88, 0.91, 1.00, 0.66],
        [-0.00, 0.76, 0.75, 0.66, 1.00],
    ]
)


# Plotting with a clean, low-ink framework
ax = sns.heatmap(
```

```

matrix_data,
annot=True,
fmt=".2f",
cmap="coolwarm",
xticklabels=labels,
yticklabels=labels,
cbar=False, # eliminate the Legend/color bar
linewidths=0.5,
)

plt.title("Spearman Correlation Matrix", fontsize=11, pad=15)
plt.tick_params(axis="both", which="both", length=0) # Removes tick mark
plt.tight_layout()
plt.show()

```



```

In [170... # isolate the two categories we want to focus on 'Just Chatting' & 'FPS'
chart_data = pd.DataFrame(
    {
        "time_of_day": ["Evening", "Evening", "Morning", "Morning"],
        "category": ["Creative", "FPS", "Creative", "FPS"],
        "avg_peak_viewers": [15663, 42797, 28500, 56000],
    }
)

# high-contrast dot plot
alt.Chart(chart_data).mark_circle(size=140).encode(
    x=alt.X(
        "time_of_day:N", title="Broadcast Window", axis=alt.Axis(labelAngle=0)
    ),
    y=alt.Y(
        "avg_peak_viewers:Q",
        title="Average Peak Viewership Count",
        axis=alt.Axis(format=",.0f", tickCount=5),
    ),
    color=alt.Color(
        "category:N",
        title="Content Category",
        scale=alt.Scale(scheme="dark2"),
    ),
)

```



```
),
).properties(
  title="Average Peak Audience Sizes: Morning vs. Evening",
  width=350,
  height=350,
).display()
```

